

Very Naive MIPS CPU using Clash

ZHU Yifan (118010469)

<i@zhuyi.fan>

May 5, 2020

In memory of Carl Quinn, for his great contributions to the field of Programming Languages.



Introduction

In CUHK(SZ) and many other universities, writing an MIPS CPU with pipeline is a required work for the architecture courses. However, most teaching materials just provide students with some basic concepts of CPUs and do not give essential introductions on languages or the potential difficulties of the implementation. Writing this book, we want to achieve the following goals:

- Give a detailed description on each part of the MIPS CPU.
- Clarify how we can write a sequential logic circuit, avoiding oscillations and other problems.
- Introduce Clash, a higher level HDL that can generate synthesizable Verilog files and reduce the complexity of development.

We are **NOT** going to implement a fully functional MIPS CPU. Instead, we will only structure some skeletons that can help us understand the concepts and principles. We hope our readers can gain some basic knowledge of hardware design and the Clash language from this book.

Contents

Introduction	ii
1 Preparation	1
1.1 Prerequisites of this Book	1
1.2 Install Icarus Verilog	1
1.3 Prepare Haskell Environment	2
2 Essential Verilog	3
2.1 Module Interface	3
2.2 Sequential Structures	5
2.2.1 Assigning Sequential Values	5
2.2.2 Test Bench	6
3 Journey to the Clash Language	10
3.1 Define Circuits	10
3.2 Useful Types	12
3.2.1 Bit and BitVector	12
3.2.2 Sized Integers	13
3.2.3 Sized Vector	13
3.2.4 Remarks on Haskell Types	14
3.3 Step into the Sequential Logic World	15
3.4 Function Utilities	16
3.4.1 ROM and RAM	16
3.4.2 Stateful Procedures	17
3.4.3 Bundle and Unbundle	20

4	Writing the CPU	22
4.1	What is CPU	22
4.2	Pipeline: Why and How	23
4.2.1	Control Hazards: Branch and Jump	23
4.2.2	Data Hazards	24
4.2.3	Special Changes	25
4.3	Implementation	26
4.3.1	Instruction Module	26
4.3.2	Decode Module	33
4.3.3	Arithmetic Module	47
4.3.4	Memory Module	57
4.3.5	Write-Back Module	62
	Advanced Topics	64
	GtkWave	67
	Sample Output	69
	Acknowledgments	71

Chapter 1

Preparation

RIGHT before our journey of implementing an MIPS CPU using Clash language, we need to get our equipment ready.

1.1 Prerequisites of this Book

Reading this book, you are expected to have some basic knowledge of Verilog HDL and Haskell. However, if you happen to have little experience on these two languages, do not worry too much; they are just the language tools that we are going to use to express the logic and thoughts. The expressions should be easy to understand and we are going to provide some detailed descriptions on those critical lines.

It is also a good idea to acquire some basic knowledge about Digital Logic Circuits. You'd better grab the concepts of clock, combinatorial logic and sequential logic.

1.2 Install Icarus Verilog

Icarus Verilog is a tool to synthesis Verilog sources and generate simulation executables. We are going to use it as our default Verilog compiler. It is available to GNU/Linux, Mac OS and Windows.

Windows users can follow this link (<http://bleyer.org/icarus/>) to download it.

Mac users can follow this tutorial

(https://blog.csdn.net/zach_z/article/details/78787509) to download it.

As for GNU/Linux users, your distributions likely provide Icarus Verilog via their repositories.

The recommended warning flags are `-Wall -Winfloop`.

1.3 Prepare Haskell Environment

We are going to use stack for the projects. It should be easy to install – just go through this document (<https://docs.haskellstack.org>) to get all the requirements settled.

As for Clash, there are several ways to install it. The binaries are ready at Nikpkg and Snapcraft. It is also doable to compile the tools from source. You are recommended to visit its website (<https://clash-lang.org>) before installing it.

After all things are settled, you should be able to play with the template project at GitHub, under dramforever/clash-with-stack (Great thanks for **dramforever**).

Dramforever also mentioned some interesting tricks to reduce code redundancy in Clash, they are not applied here, see the Advanced Topics at our appendix or visit the repository.

This template does not use `mtl` library that we need (for the state monad), you may need to add it by yourself to the `package.yaml`.

Chapter 2

Essential Verilog

WE would like to introduce some basic Verilog knowledge that will be used in this book. We will not go into details here, just showing some most common cases. Let us first take a look at the final outcome of our project.

2.1 Module Interface

```
1  `timescale 100fs/100fs
2  module CPU
3      ( input  CLOCK // clock
4        , input  RESET // reset
5        , input  ENABLE
6      );
7      wire [32:0] BRANCH;
8      wire [30:0] PC_INSTRUCTION;
9      wire [31:0] PC_VALUE;
10     wire [37:0] WRITE_PAIR;
11     wire          STALL;
12     wire [5:0]    DM_WRITE;
13     wire [1:0]    DM_MEM;
14     //...
15
```



```

16     InstructionModule IM
17     ( // Inputs
18       .CLOCK(CLOCK), // clock
19       .RESET(RESET), // reset
20       .ENABLE(ENABLE),
21       .BRANCH(BRANCH),
22       .STALL(STALL),
23
24       // Outputs
25       .PC_INSTRUCTION(PC_INSTRUCTION),
26       .PC_VALUE(PC_VALUE)
27     );
28     // ...
29
30     assign AM_FW_0 = MMO_WRITE_PAIR;
31
32     assign AM_FW_1 = WB_WRITE_PAIR;
33
34     assign WRITE_PAIR = WB_WRITE_PAIR;
35
36     assign BRANCH = WB_BRANCH;
37
38     endmodule

```

These lines are extracted from the CPU source code.

1. The first line is a compiler derivative that defines the precision of timing;
2. Line 2 to line 6 define the module interface of a CPU, with three input ports. As for outputs, you can add something like **output wire** OUTPUT;
3. Line 7 to line 13 declare some wire variables. As you can see, the number of bits in the declaration can be customized. Wires are used to connect different components of the circuits, you can treat them as alias of the original port because the value of a wire is refreshed as soon as its input side changes.

4. Apart from wires, another commonly used thing is register (**reg** [31:0] REGISTER); you can treat them as the variables in the common sense, which has its own state.
5. `InstructionModule IM (...)` declares a component named IM whose definition is in another module called `InstructionModule`;
6. There are mainly two ways to interact with another module, one is to use it as what is shown in the code: use a syntax like `.PORT(SOME_WIRE)` to connect the inputs and outputs; the other way is commonly used in debugging: you can get the value of the components in another module via something like `IM.CLOCK`;
7. Those **assign** statements are used to connect the wires.

2.2 Sequential Structures

The previous example is just about connecting wires, however, many circuits also need to handle sequential events.

2.2.1 Assigning Sequential Values

There are several ways to assigning values in a sequential logic environment:

```
1 module example();  
2     reg A, B, C;  
3     initial begin  
4         A = 1;  
5         B = A;  
6         C = B;  
7     end  
8 endmodule
```

This example shows a way to assign the initial values of some registers within an **initial** block; Notice that operator `=` stands for the blocking assignment, which means the assignments will happen one by one.

What if we want to handle the assignment at some specific time? We can then use a statement in the form of **always** `@(... sensitivity list ...) begin,`

the following example shows the non-blocking assignments happening on each rising edge of the clock:

```
1  always @(posedge CLOCK) begin
2      B <= A;
3      C <= B;
4      D <= C;
5  end
```

If you want to describe a combinatorial logic, you should use `always@(*)` (only blocking assignment should be used within the scope, otherwise it is likely to generate unexpected oscillations – the circuit will never reach a stable state). The event within the block will be triggered as long as any of the inputs changes.

2.2.2 Test Bench

```
1  module TEST();
2      reg clk, reset, enable;
3      initial
4          begin
5              clk = 0;
6              reset = 0;
7              enable = 1;
8          end
9
10     always
11         #1000 clk = !clk;
12
13
14     CPU cpu(clk, reset, enable);
15
16
17
18     initial begin
```

```

19 $monitor(
20 "=====\\n",
21 "TIME:          %-d\\n",          $time,
22 "STALL:         %b\\n",          cpu.STALL,
23 "----- Instruction -----\\n",
24 "PC/4 + 1:      %-d\\n",          cpu.IM.PC_VALUE,
25 "INSTRUCTION:   %b\\n",          cpu.IM.result_1[31:0],
26 "INSTRUCTION:   %b [inner form]\\n", cpu.IM.PC_INSTRUCTION,
27 "----- Decode -----\\n",
28 "RS:            %-d\\n",          cpu.DM_RS,
29 "RS VALUE:      %b\\n",          cpu.DM_RSV,
30 "RT:            %-d\\n",          cpu.DM_RT,
31 "RT VALUE:      %b\\n",          cpu.DM_RTV,
32 "MEM_OP:        %b\\n",          cpu.DM_MEM,
33 "REG_WRITE:     %b\\n",          cpu.DM_WRITE,
34 "ALU_CTL:       %b\\n",          cpu.DM_ALU,
35 "IMMEDIATE:     %b\\n",          cpu.DM_IMM,
36 "STAGE_PC/4 + 1: %-d\\n",          cpu.DM_COUNTER,
37 "----- Arithmetic -----\\n",
38 "REG_WRITE:     %b\\n",          cpu.AM_WRITE_REG,
39 "MEM_OP:        %b\\n",          cpu.AM_MEM_OP,
40 "ALU_RESULT:    %b\\n",          cpu.AM_RESULT,
41 "BRANCH_TARGET: %b\\n",          cpu.AM_BRANCH_TARGET,
42 "----- Memory -----\\n",
43 "BRANCH_TARGET: %b\\n",          cpu.MMO_BRANCH,
44 "WRITE_BACK:    %b\\n",          cpu.MMO_WRITE_PAIR,
45 "NEXT_FETCH_ADDRESS: %-x\\n",
    ⇨ cpu.MM.MainMemory_res.FETCH_ADDRESS,
46 "FETCH_RESULT:  %-x\\n",          cpu.MM.MainMemory_res.DATA,
47 "WRITE_SERIAL:  %b\\n",
    ⇨ cpu.MM.MainMemory_res.EDIT_SERIAL,
48 "----- Write Back -----\\n",

```

```

49 "BRANCH_TARGET:      %b\n",      cpu.WB_BRANCH,
50 "WRITE_BACK:        %b\n",      cpu.WB_WRITE_PAIR,
51 "===== \n"
52 );
53     #100000 $finish();
54     end
55
56 endmodule

```

Here is the test bench that we are going to use. Those #XXXX statements mean delaying for the given amount of time before the event. Hence,

always

```
#1000 clk = !clk;
```

actually defines a clock with period 2000. There are several special functions we are going to use,

- **\$finish()** terminates the simulation
- **\$stop()** pauses the simulation
- **\$display("format string", a, "format string", b)** displays the instant value of the variables; basic formats are:
 - %d: digits
 - %-d: digits (left aligned)
 - %b: binary
 - %x: hexadecimal
- **\$monitor("format string", a, "format string", b)** used the same as display, but it will be triggered everytime a monitored variable updates
- **\$time** gets the current time
- **\$readmemb("file.bin", BLOCK)**; initializes a large memory block with a file

- `$dumpfile("file.vcd")` dumps IEEE standard vcd files. These files can be visualized by some softwares like GtkWave to provide a handy way to debug.
- `$dumpvars(0, cpu)` sets the value and module to dump; level 0 will automatically dumps the variables in the module recursively while level 1 will only dumps those manually listed variables.

Chapter 3

Journey to the Clash Language

CLASH will be our main language to write the CPU. Clash supports most of Haskell syntax, yet it cannot support some advanced features such as GADT pattern matching. To see the full list of limitations, please check its official tutorial (<http://hackage.haskell.org/package/clash-prelude-1.2.0/docs/Clash-Tutorial.html>). It will also be a great idea to go through the troubleshooting part if you face some difficulties later.

3.1 Define Circuits

You can simply write a circuit in the way of writing a Haskell function:

```
module Example where
orGate :: Bool -> Bool -> Bool
orGate = (||)
```

How to generate a verilog module from the code? If your function is named as `topEntity`, just load the `clash.clashi` on your own or using `stack` and then input `:verilog Example` in the REPL, then the outputs are ready at the `verilog` subdirectory under your working directory.

However, in most cases, you need to write a special annotation for the function:

```
{-# ANN orGate
  ( Synthesize
```

```

    { t_name = "OrGate"
    , t_inputs =
        [ PortName "X"
        , PortName "Y"
        ]
    , t_output =
        PortName "RESULT"
    }
)
#-}

```

You can customize the names of the ports or the whole module. There is another fantastic thing that you can also set a test bench for your circuits. As we are not going to use Clash to generate test benches, it is up to you to investigate. Here is the link:

<http://hackage.haskell.org/package/clash-prelude-1.2.1/docs/Clash-Annotations-TopEntity.html#v:TestBench>

Here is another example to demonstrate how to handle product ports.

```

{-# ANN someGates
(Synthesize
  { t_name = "SomeGates"
  , t_inputs =
      [ PortName "X"
      , PortProduct "IN"
          [ PortName "Y"
          , PortName "Z"
          ]
      ]
  , t_output = PortProduct "OUT"
      [ PortName "0"
      , PortName "1"
      , PortName "2"
      ]
  }
)

```



```

    }
  )
#-}

```

This annotation can be used with functions in the form of

```
someGate :: x -> (y, z) -> (o0, o1, o2)
```

3.2 Useful Types

3.2.1 Bit and BitVector

Bit is just bit and BitVector is just a statically sized vector of bits. As a single Bit and Bool are quite the same, **Clash.Prelude** provides some handy functions for us to convert them from each other.

```

boolToBit :: Bool -> Bit
bitToBool :: Bit -> Bool
boolToBV :: KnownNat n => Bool -> BitVector (n + 1)

```

BitVector can be sliced and indexed, the following code shows some examples:

```

(!) :: (BitPack a, Enum i) => a -> i -> Bit
-- ^ BitVector is within the class of BitPack, so we can use (!)
--   operator to fetch the value

vec ! 0
-- ^ fetch the lowest bit

slice :: (BitPack a, BitSize a ~ ((m + 1) + i))
      => SNat m -> SNat n -> a -> BitVector ((m + 1) - n)
-- ^ fetch a slice from the the bit vector
-- ^ because the vector is defined as a dependent type
-- ^ this interface is a bit terrifying, but it is handy to use

```

```
slice d31 d20 vec
-- ^ slice from index 31 to index 20
```

As you can see, `d0` to `d1024` are predefined literals for static natural numbers, you can use them for the vector index.

How about the bitwise operations? There is a whole set of operators and functions.

```
(.|.) :: Bits a => a -> a -> a      -- ^ bitwise or
(&.) :: Bits a => a -> a -> a      -- ^ bitwise and
xor :: Bits a => a -> a -> a       -- ^ bitwise xor
complement :: Bits a => a -> a     -- ^ bitwise not
shiftR :: Bits a => a -> Int -> a   -- ^ bitwise shift
shiftL :: Bits a => a -> Int -> a   -- ^ bitwise shift
unsafeShiftR :: Bits a => a -> Int -> a -- ^ bitwise shift
```

3.2.2 Sized Integers

There are mainly two types of sized integers: **Unsigned** (`n :: Nat`) and **Signed** (`n :: Nat`). Most bitwise operations can also be applied to sized integers, but please notice that for right shifting, the normal version takes care of the sign bit and the unsafe version just do the logical shifting.

It is also possible to extend sized integers and BitVector,

```
extend :: (Resize f, KnownNat a, KnownNat b) => f a -> f (b + a)
```

It is very handy that extensions will handle the changing of the sign bits automatically.

Although BitVector and sized integers are very similar, you cannot treat them as the same thing. However, if you need to convert the types, you can use the following functions:

```
pack :: BitPack a => a -> BitVector (BitSize a)
unpack :: BitPack a => BitVector (BitSize a) -> a
```

3.2.3 Sized Vector

Another important type is sized vector: **Vec** `:: Nat -> * -> *`. In fact, the memory blocks that we are going to use are just wrappers around this type (`Vec 512 (BitVector 32)`).

To get the value at a specific index, you can use the following function:

```
(!!) :: (KnownNat n, Enum i) => Vec n a -> i -> a
```

To update the data, you can use this function:

```
replace :: (KnownNat n, Enum i) => i -> a -> Vec n a -> Vec n a
```

Notice that, consecutive update at the single clock (which is really a bad decision) may make Clash fail to synthesize the circuit. Sometimes we need to provide an initial value for the RAM, then the following function will be handy:

```
replicate :: SNat n -> a -> Vec n a
```

```
replicate d512 0 :: Vec 512 (BitVector 32)
-- ^ create zero-filled vector
```

3.2.4 Remarks on Haskell Types

Other Haskell types like **Maybe** can also be used without problem. For example, **Maybe Bool** will be represented by 2 bits in Verilog.

```
Just True  = 11
Just False = 10
Nothong    = 0x
```

Clash will also find a way to represent your own defined sum types, for example,

```
data MyEnum = A | B | C | D
```

will be represented by something like

```
A -> 00
B -> 01
C -> 10
D -> 11
```

3.3 Step into the Sequential Logic World

The previous parts are mainly talking about combinatorial logic; how about the sequential one?

In Clash, sequential logic things are wrapped into a type **Signal** (`dom :: Symbol`) a. The `dom` stands for the signal domain, which provides some basic configurations such as clock, reset, enable, frequency and etc. The default domain is **System**, which is the standard global domain. It is also possible to define domains on your own and setup some multiple clock domains; these advanced features are not used in this book.

Signal is not Monad, but it provides the interfaces of Functor and Applicative.

To check the content of Signal, you can use a function called `sampleN`, to sample it.

To test your code within Signal environment, you can also use `enableGen` to generate enable signals, `clockGen` to generate clock signals and `resetGen` to generate reset signals. What's more, there is also a `simulate` function, which allows you to use a syntax like

```
simulate @System myInterface
```

to generate the result.

Usually, enable, clock and reset signals are required everywhere within a synchronous circuit. Imagine writing these three signals repeatedly at every function, it will definitely become tedious. Hence, Clash provides a special way to define a generalized signal domain which hides some global signals, you can simply expose them at those interfaces to synthesize. The following example illustrates how to use this feature

```

1 example :: HiddenClockResetEnable dom
2   => Signal dom Bool
3   -> Signal dom Bool
4 example = fmap complement
5
6 example'
7   :: Clock System
8   -> Reset System
9   -> Enable System
```

```

10     -> Signal System Bool
11     -> Signal System Bool
12 example' = exposeClockResetEnable example

```

3.4 Function Utilities

3.4.1 ROM and RAM

Clash provides some predefined functions for us to declare large blocks of memory.

Asynchronous Memory

Let us first have a look at the asynchronous ROM and RAM,

```

asyncRom :: (KnownNat n, Enum addr)
=> Vec n a -- ^ initial vector
-> addr    -- ^ read address
-> a       -- ^ read result

asyncRam
:: ( Enum addr, KnownDomain dom
    , GHC.Classes.IP (AppendSymbol dom "_clk") (Clock dom)
    , GHC.Classes.IP (AppendSymbol dom "_en") (Enable dom)) => SNat n
-> Signal dom addr -- ^ read address
-> Signal dom (Maybe (addr, a)) -- ^ write data
-> Signal dom a -- ^ read result

```

The asynchronous version will output the content in the read address at the same clock cycle. If the read address and the write address conflicts, the default strategy is write-after-read; however, you can use `readNew . asyncRam` to apply the read-after-write strategy.

Synchronous Memory

Asynchronous memory is handy enough, but it is not the optimal structure: asynchronous memory may require a lot of LUTs in FPGA and the cost must be consid-

ered if the memory size is relatively large. Fortunately, there is a synchronous version of RAM, it corresponds to the BRAM structure in FPGA. However, there is a big difference that the read and write operation issued in the current clock cycle will generate outcome in the next cycle; we must take care of this feature when designing circuits.

blockRam

```

:: ( KnownDomain dom
    , GHC.Classes.IP (AppendSymbol dom "_clk") (Clock dom)
    , GHC.Classes.IP (AppendSymbol dom "_en") (Enable dom), NFDDataX a
    , Enum addr )
=> Vec n a
-> Signal dom addr
-> Signal dom (Maybe (addr, a))
-> Signal dom a

```

Similarly, you can change the read-write conflict resolution.

For asynchronous ROM and synchronous RAM, Clash also provides some functions like `blockRamFile`, which will be translated into some Verilog code using `readmemb` function; which makes it possible for us to initialize the memory field using external files.

3.4.2 Stateful Procedures

There are several ways to handle stateful procedures.

Register

Clash provides a `register` function to help us to handle some basic stateful procedures. Register buffers and delays the input signal by one clock tick before output, and in the first clock tick the output is the initial value.

register

```

:: ( HiddenClockResetEnable dom
    , NFDDataX a )
=> a

```

```

-> Signal dom a
-> Signal dom a

```

```

register 1 -- ^ declare a register with initial value 1

```

Mealy

Mealy Machine is a sort of state machine whose output is determined by the current input and state.

```

mealy
  :: ( HiddenClockResetEnable dom
      , NFDataX s )
  => (s -> i -> (s, o))
  -> s
  -> Signal dom i
  -> Signal dom o

```

Let us write a special counter using Mealy Machine: if the outside provides an input, it will set it as the next counting value, otherwise, it just increases the counter and outputs the current value.

It seems that we can describe the state transformation with the following function:

```

counterT
  :: Unsigned 32
  -> Maybe (Unsigned 32)
  -> (Unsigned 32, Unsigned 32)
counterT state Nothing = (state + 1, state)
counterT state (Just x) = (x      , state)

```

To transform `counterT` into a state machine, just apply the `mealy` function together with an initial value to it

```

counter
  :: HiddenClockResetEnable dom
  => Signal dom (Maybe (Unsigned 32))

```

```

=> Signal dom (Unsigned 32)
counter = mealy counterT 0

```

Moore

Another kind of commonly used finite state machine is the Moore Machine, whose output is determined only by the current state.

```

moore
  :: ( HiddenClockResetEnable dom
      , NFDataX s )
  => (s -> i -> s)
  -> (s -> o)
  -> s
  -> (Signal dom i -> Signal dom o)

```

The counter example can also be transformed into a Moore Machine:

```

counterT
  :: Unsigned 32
  -> Maybe (Unsigned 32)
  -> Unsigned 32
counterT state Nothing = state + 1
counterT _      (Just x) = x

```

```

counter
  :: HiddenClockResetEnable dom
  => Signal dom (Maybe (Unsigned 32))
  -> Signal dom (Unsigned 32)
counter = moore counterT id 0

```

State Monad

It is also possible to use State Monad to implement the state machine. Still use the counter as an example:


```

counterS
  :: Maybe (Unsigned 32)
  -> State (Unsigned 32) (Unsigned 32)
counterS input = do
  state <- get
  put $ maybe (state + 1) id input
  return state

```

However, we need to write a function to transform our state monad into a synthesizable state machine:

```

asStateM
  :: (HiddenClockResetEnable dom, NFDDataX s)
  => (i -> State s o)
  -> s
  -> (Signal dom i -> Signal dom o)
asStateM f i = mealy g i
  where
    g s x =
      let (o, s') = runState (f x) s
      in (s', o)

```

With this function, we are able to write something like:

```
asStateM counterS 0
```

3.4.3 Bundle and Unbundle

Consider we have the following functions

```

foo
  :: HiddenClockResetEnable dom
  => Signal dom Input
  -> Signal dom (TypeA, TypeB)

bar

```

```
:: HiddenClockResetEnable dom
=> Signal dom TypeA
-> Signal dom TypeB
-> Signal dom Result
```

The problem emerges when we want to combine these two functions: it is hard to take out the value of each part of the tuple wrapped in the signal environment.

Hence, we can use the `unbundle` function.

```
let (a, b) = unbundle $ foo input
in bar a b
```

The `bundle` function just acts as the inverse of `unbundle`:

```
a :: Signal System A
b :: Signal System B
func :: Signal System (A, B) -> Signal System C
res :: Signal System C
res = func $ bundle (a, b)
```

Chapter 4

Writing the CPU

NOW, we are fully armed with our equipment. It is the time to get our hands dirty and start implementing a very naive MIPS CPU.

4.1 What is CPU

Central Process Unit, a.k.a. CPU, is the most important part of computers. It is in charge of the memory loading and storing, arithmetic operations and lots of other important stuffs. Most CPUs consist of a register called program counter, which record the current instruction position in the memory; in each cycle, the CPU fetches an instruction according to the program counter and start to handling a series of events encoded in the instruction.

The CPU we are going to implement consists of five parts:

1. Instruction Module: Fetch the instruction and maintain the value of the program counter.
2. Decode Module: Decode the instruction, determine the operations to be executed and get the value of the operands from the register file.
3. Arithmetic Unit: Execute the arithmetic operations, determine the branch targets and memory operations.
4. Memory Module: Load and store data from and into memory.

5. Write Back Module: Act as a transition module before register writing and branching.

4.2 Pipeline: Why and How

Although the functions of the CPU can be achieved within a single cycle: on each rising edge of the clock, we just fetch a new instruction and wait until all the required operations are finished, it is apparent that the cycle may become too long and inefficient.

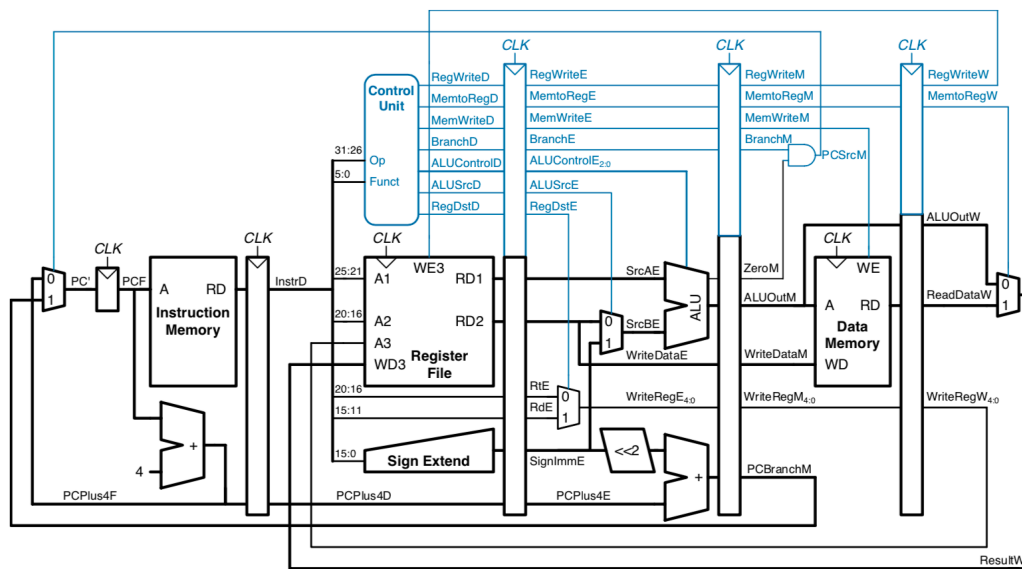


Figure 4.1: Classical MIPS pipelines (without any hazard handling)

Hence, many CPUs use a strategy called pipeline which split the execution into multiple stages. Each stage stores the previous input (the output from the predecessor stage) in its register, and when the clock rises, the CPU is handling multiple instructions with different facilities. The efficiency is thus largely improved.

However, lots of new problems emerge because of the pipeline.

4.2.1 Control Hazards: Branch and Jump

Just consider the CPU shown in Figure 4.1. Without pipeline, a branch instruction triggers the branching in the single cycle and setup the new PC value within the same

cycle; the next instruction to execute will always be correct. However, in the pipeline model, if the branch condition fails, everything works fine; nevertheless, when the branch condition checks and the branch instruction reaches the Write-Back stage, there will be three invalid instructions already accumulated in the pipeline. We must find a way to flush those invalid instruction.

This is actually easy: we just check the branch signal; if the branch target is set, we send a signal to all four modules before Write-Back and ask them to clear their instructions. **Notice that the memory writing should be stopped immediately while the register writing should be allowed**, because there are jumping instructions like `jal` that will store the current PC value into a register and this operation is always valid as it arrives at the Write-Back part at the same time of the branch target.

In most CPUs, the cost of branching can be further reduced by branch prediction and moving branching checking into an earlier stage. Here, we are not going to care about these strategies and just handle the control hazards with stalling.

4.2.2 Data Hazards

There are several kinds of data hazards:

1. Consecutive arithmetic operations with data dependency:

```

1  addi $a0, $a0, 1
2  addi $a0, $a0, 1
3  addi $a0, $a0, 1

```

Notice that when the second/third instruction reaches the Arithmetic Module, the first/second instruction is at the Memory Part/Write-Back Part; which means the write-back operation is not applied. Hence, we must figure out a way to forward the result stored in Memory Part/Write-Back Part to Arithmetic Module in advance.

2. Arithmetic Operation followed by SW:

```

1  addi $a0, $a0, 1
2  sw   $a0, -4($sp)

```

In the SW case, the value of register a0 is not ready. We can simplify the handling by moving determination of write value into Arithmetic Part. Hence, using the strategy of forwarding, we can also handle this problem.

3. LW followed by Arithmetic Operation

```
1  lw    $t0, 0($sp)
2  addi  $t0, $t0, 1
```

In the pipeline shown in Figure 4.1, the load value will not be ready even with forwarding. In this case, we must stall the pipeline: keep the PC value, the state of Decode Module and the Arithmetic Module unchanged (but accepting register write-back) and insert a NOP state into memory module for the next cycle while handling the current operation of loading. Therefore, in the next cycle, the second instruction is still at the Arithmetic Module while the load result reaches the Write-Back Part which makes it possible to be forwarded to ALU.

4.2.3 Special Changes

We made several special changes to make life easier:

1. Because we are using the BRAM structure, which means that there is one cycle delay before we can get the real data, we need to forward the ALU results to the Memory Module in the same cycle so that in the next cycle the Memory Output is exactly what the instruction in the Memory Module requires.

```
1  j      SOME_PLACE
2  sw     $zero, 4($sp)
3  sw     $zero, 4($sp)
```

Previously, when the branch target reaches Write-Back Part, the Memory Operation hasn't taken place; however, in our case, the stall caused by branching will only prevent the second Memory Operation after branching; we need to also check the branching before we want to write something into the memory. Fortunately, this is always available since, when the instruction right after branching arrives ALU, the branching instruction is exactly at the Memory Part, we can simply check whether the branching target is set or not.

2. There is no need to care about the Load Data Hazards as we will also forward the data together with the memory fetching result. (Thanks to BRAM delay, the extra cost is little)

4.3 Implementation

4.3.1 Instruction Module

Instruction Set

We will only implement a very small set of MIPS instructions. Let's first write some functions to help us decode the instruction:

```
-- The Format of MIPS Instructions
data Format
  = NoType          -- NoType (Specialized for NOP Instruction)
  | RType           -- R-Type Instruction Format
    (BitVector 6)   -- Operation Code
    (BitVector 5)   -- Register S
    (BitVector 5)   -- Register T
    (BitVector 5)   -- Register D
    (BitVector 5)   -- Extra Information for Shifting Amount
    (BitVector 6)   -- Function Code
  | IType           -- I-Type Instruction Format
    (BitVector 6)   -- Operation Code
    (BitVector 5)   -- Register S
    (BitVector 5)   -- Register T
    (BitVector 16)  -- Immediate Value
  | JType           -- J-Type Instruction Format
    (BitVector 6)   -- Operation Code
    (BitVector 26)  -- Jump Target
deriving Show
```

We will first recognize the instruction format and then transform it into each recognized instruction.

```

decodeFormat :: BitVector 32 -> Format
decodeFormat 0 = NoType
decodeFormat vec =
    let opcode = slice d31 d26 vec
    in case opcode of
        0 ->
            RType
                <$> (slice d31 d26)
                <*> (slice d25 d21)
                <*> (slice d20 d16)
                <*> (slice d15 d11)
                <*> (slice d10 d6)
                <*> (slice d5 d0)
                $ vec
        code ->
            if code == 0b000010 || code == 0b000011
            then JType
                <$> (slice d31 d26)
                <*> (slice d25 d0)
                $ vec
            else IType
                <$> (slice d31 d26)
                <*> (slice d25 d21)
                <*> (slice d20 d16)
                <*> (slice d15 d0)
                $ vec

```

The format decode is trivial:

- An all-zero instruction is an NOP;
- Otherwise, instructions with 0 opcode will be dispatched into R-Format;
- Instructions with special jumping opcode will be dispatched into J-Format;

- Other instructions are in the I-Format

After the format is determined, we can then transform instructions into our inner forms:

```

type Register = Unsigned 5
data Instruction
  = NOP
  | ADD Register Register Register
  | ADDI Register Register (Signed 16)
  | ADDU Register Register Register
  | ADDIU Register Register (Unsigned 16)
  | SUB Register Register Register
  | SUBU Register Register Register
  | AND Register Register Register
  | ANDI Register Register (BitVector 16)
  | NOR Register Register Register
  | OR Register Register Register
  | ORI Register Register (BitVector 16)
  | XOR Register Register Register
  | XORI Register Register (BitVector 16)
  | BEQ Register Register (Signed 16)
  | BNE Register Register (Signed 16)
  | SLT Register Register Register
  | SLTI Register Register (Signed 16)
  | SLTU Register Register Register
  | SLTIU Register Register (Unsigned 16)
  | LW Register Register (Signed 16)
  | SW Register Register (Signed 16)
  | SLL Register Register (Unsigned 5)
  | SRL Register Register (Unsigned 5)
  | SRA Register Register (Unsigned 5)
  | SLLV Register Register Register
  | SRLV Register Register Register
  | SRAV Register Register Register

```

```

| J (Unsigned 26)
| JAL (Unsigned 26)
| JR (Unsigned 5)
deriving Show
deriving Generic
deriving NFDataX

```

We do not provide the function

```
decodeTyped :: Format -> Instruction
```

here as it should be easy to write and just require some repeated works to recognize the instruction based on the opcode and the function code.

Just mention a small trick: to reduce code repetition, you can use the Applicative property of Readers):

```

t1 (x, _, _) = x
t2 (_, y, _) = y
t3 (_, _, z) = z
makeType func = func <$> unpack . t1 <*> unpack . t2 <*> unpack . t3
-- then you can just write something like:
case fn of
  0b100000 -> (makeType ADD) (rs, rt, rd)
  0b100001 -> (makeType ADDU) (rs, rt, rd)
  0b100100 -> (makeType AND) (rs, rt, rd)
  -- more cases ...

```

Please notice that for most R instructions, we will keep the format as **ADD** rs rt rd, but for those shift related operations, the format will be **SLL** rd rt sa; those I instructions adapts the format of **ADDI** rs rt imm.

The Instruction RAM

We will use BRAM to implement the instruction memory space. The memory itself is quite simple, we just use the `blockRamFile` function to achieve the goal

```

instrRAM'
  :: HiddenClockResetEnable dom
  => Signal dom MemAddr
  -> Signal dom (BitVector 32)
instrRAM' = (flip $ blockRamFile d512 "instructions.bin") $ pure
  ↪ Nothing

instrRAM
  :: Clock System
  -> Reset System
  -> Enable System
  -> Signal System MemAddr
  -> Signal System (BitVector 32)
instrRAM = exposeClockResetEnable instrRAM'

```

Instruction Module Interface

Now we must think carefully of what we are going to in this module:

1. We need to maintain the value of program counter, if no branching happens, we just increase it; otherwise, we set the counter to the branch target immediately.
2. On Each cycle, we will need to output two things: an instruction and (PC + 1) (one plus the instruction index). As BRAM has a delay effect, the current PC in this cycle is exactly the value we want, hence we do not need to make an extra addition to the counter.
3. When stall happens, output an NOP. As the branch target is set immediately, there is no need to keep the old PC value.

With these ideas, we can come out the following design:

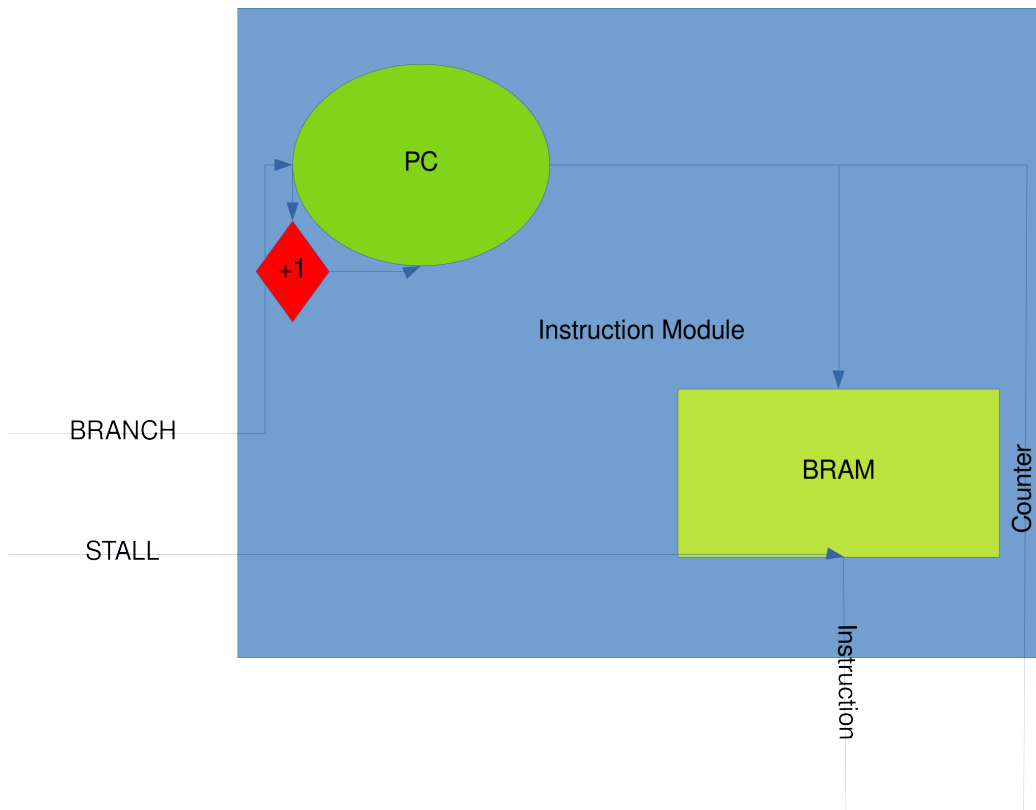


Figure 4.2: Instruction Module Diagram

Let's first maintain the PC value using a Mealy Machine:

```

type PCInput = Maybe (Unsigned 32)
programCounterT
  :: Unsigned 32
  -> PCInput
  -> (Unsigned 32, Unsigned 32)
programCounterT state (Just t) = (t + 1, t)
programCounterT state _ = (state + 1, state)

programCounter
  :: HiddenClockResetEnable dom
  => Signal dom PCInput
  -> Signal dom (Unsigned 32)

```

```
programCounter = mealy programCounterT 0
```

As you can see, unless the branch target is set, the PC just behaves as a monotonic counter.

```

1  {-# ANN pcModule
2      (Synthesize{t_name = "InstructionModule",
3          t_inputs =
4              [PortName "CLOCK", PortName "RESET", PortName "ENABLE",
↪   PortName "STALL", PortName "BRANCH"],
5          t_output =
6              PortProduct "PC" [PortName "INSTRUCTION", PortName
↪   "VALUE"]})
7      #-}
8
9  pcModule
10     :: Clock System
11     -> Reset System
12     -> Enable System
13     -> Signal System Bool
14     -> Signal System (Maybe (Unsigned 32))
15     -> Signal System (Instruction, (Unsigned 32))
16  pcModule clk rst enable stall br = bundle (instr, next)
17     where
18         programCounter' = (exposeClockResetEnable programCounter) clk
↪         rst enable
19         next = programCounter' $ br
20         ram = instrRAM clk rst enable next
21         ram' = decodeTyped . decodeFormat <$> ram
22         instr' op ram =
23             case op of
24                 False -> ram
25                 True -> NOP
26         instr = instr' <$> stall <*> ram'

```

If you are not familiar with Haskell, just notice that `< $ >`, `< * >` are the operators using the Applicative property of `Signal` to lift a function of `f :: a -> b -> c` type onto `f :: Signal dom a -> Signal dom b -> Signal dom c`, so that we can transform combinatorial logic into a sequential form. The data flow is simple:

1. Get the current counter with branch flag
2. Get the memory output and set the next address to fetch
3. Decode the output
4. Check the stall condition and output the instruction and PC value

Now we can happily generate the Verilog code for this module, try it!

4.3.2 Decode Module

Register File

Decode Module contains the Register File which is in charge of maintaining the register value. Let us first take a look at these definitions. As the logic relation is a little bit complicated, we can use State Monad to construct the state machine for the Register File. Notice that we require `LambdaCase` extension in some of the following code.

```

1  type Reg = BitVector 32
2
3  type RegNo = Unsigned 5
4
5  registerFileS
6    :: ( RegNo -- rs
7        , RegNo -- rt
8        , Maybe (RegNo, Reg) -- write register
9        )
10     -> State (Vec 32 Reg) (Reg, Reg)
11  registerFileS (reg0, reg1, writePair) = do
12    regs <- get
13    let newS =

```

```

14         case writePair of
15             Nothing -> regs
16             Just (a, b) ->
17                 if a /= 0
18                     then replace a b regs
19                     else regs
20     res0 = newS !! reg0
21     res1 = newS !! reg1
22     put newS
23     return (res0, res1)
24
25 {-# ANN registerFile
26     (Synthesize
27         { t_name = "RegisterFile"
28         , t_inputs =
29             [ PortName "CLOCK"
30             , PortName "RESET"
31             , PortName "ENABLE"
32             , PortProduct "RF"
33               [ PortName "RS"
34               , PortName "RT"
35               , PortName "WRITE"
36               ]
37         ]
38         , t_output =
39             PortProduct "RF"
40               [ PortName "RSV"
41               , PortName "RTV"
42               ]
43         }
44     )
45     #-}

```

```

46
47 registerFile
48   :: Clock System
49   -> Reset System
50   -> Enable System
51   -> Signal System (RegNo, RegNo, Maybe (RegNo, Reg))
52   -> Signal System (Reg, Reg)
53 registerFile = exposeClockResetEnable $ asStateM registerFileS
    ↪ (replicate d32 0)

```

The Register File takes three input: two register indices to fetch and one register write serial consists of register number and target value. The outputs are the values of required registers. Here we use read-after-write strategy to make sure that the fetched data is the latest.

Control Unit

Now comes one of the most tedious part of the entire CPU, the Control Unit. It extract the information from the instruction and separate them for further execution.

The extraction is not stateful, to reduce the complexity, we can argue this part in the combinatorial form and lift it later.

Here we are going to define lots of data format for later uses:

- **MemoryOperation**: this data tags the operation on the RAM: none, read or write:

```

data MemoryOperation
  = MemNone
  | MemLoad
  | MemWrite
deriving (Generic)
deriving (NFDataX)
deriving (Show)

```

- **BranchFlag**: Four types: no-branching, branch-on-equal, branch-on-different,

jump; because conditional branching uses up two operands, we need additional field to carry the branching difference.

```
data BranchFlag
  = NoBranch
  | BranchEQ (BitVector 32)
  | BranchNE (BitVector 32)
  | Jump
  deriving (Generic)
  deriving (NFDataX)
```

- **ALUOperation**: Defines the ALU Operations, those add, sub, set, right shift operations are accompanied with an extra bit to tag whether the data is signed or not.

```
data ALUOperation
  = ALUAdd Bool
  | ALUSub Bool
  | ALUAnd
  | ALUNor
  | ALUOr
  | ALUXor
  | ALUSet Bool
  | ALUShiftL
  | ALUShiftR Bool
  | ALUNone
  deriving (Generic)
  deriving (NFDataX)
```

Now, we can write some functions to dispatch these flags. First, we can check whether there is a need of register writing. The correct register to write may vary with different instructions. One special case is the JAL instruction, which requires the update of register ra.

```
writeRegister
  :: HiddenClockResetEnable dom
```

```

=> Signal dom Instruction
-> Signal dom (Maybe (Unsigned 5))
writeRegister = fmap $
  \case
    ADD  _ _ rd -> Just rd
    ADDI _ rt _ -> Just rt
    ADDU _ _ rd -> Just rd
    ADDIU _ rt _ -> Just rt
    SUB  _ _ rd -> Just rd
    SUBU _ _ rd -> Just rd
    AND  _ _ rd -> Just rd
    ANDI _ rt _ -> Just rt
    NOR  _ _ rd -> Just rd
    OR   _ _ rd -> Just rd
    ORI  _ rt _ -> Just rt
    XOR  _ _ rd -> Just rd
    XORI _ rt _ -> Just rt
    SLT  _ _ rd -> Just rd
    SLTI _ rt _ -> Just rt
    SLTU _ _ rd -> Just rd
    SLTIU _ rt _ -> Just rt
    SLL  rd _ _ -> Just rd
    SRL  rd _ _ -> Just rd
    SRA  rd _ _ -> Just rd
    SLLV _ _ rd -> Just rd
    SRLV _ _ rd -> Just rd
    SRAV _ _ rd -> Just rd
    LW   _ rt _ -> Just rt
    JAL  _      -> Just 31
    _    _      -> Nothing

```

Then, we can also check the requirement of memory operation, branching and immediate value. Here we must pay attention to the order of extend and pack to make sure

that sign information is kept correctly.

memoryOperation

```

:: HiddenClockResetEnable dom
=> Signal dom Instruction
-> Signal dom MemoryOperation

```

memoryOperation = fmap \$

```

\case
  LW _ _ _ -> MemLoad
  SW _ _ _ -> MemWrite
  _      -> MemNone

```

branchFlag

```

:: HiddenClockResetEnable dom
=> Signal dom Instruction
-> Signal dom BranchFlag

```

branchFlag = fmap \$

```

\case
  BEQ _ _ x -> BranchEQ (pack $ extend x)
  BNE _ _ x -> BranchNE (pack $ extend x)
  JR  _      -> Jump
  J   _      -> Jump
  JAL _      -> Jump
  _   _      -> NoBranch

```

immediateValue

```

:: HiddenClockResetEnable dom
=> Signal dom Instruction
-> Signal dom (Maybe (BitVector 32))

```

immediateValue = fmap \$

```

\case
  ADDI _ _ x -> Just (pack $ extend x)
  ADDIU _ _ x -> Just (pack $ extend x)

```

```

ANDI  _ _ x -> Just (extend x)
ORI   _ _ x -> Just (extend x)
XORI  _ _ x -> Just (extend x)
SLTI  _ _ x -> Just (pack $ extend x)
SLTIU _ _ x -> Just (pack $ extend x)
LW    _ _ x -> Just (pack $ extend x)
SW    _ _ x -> Just (pack $ extend x)
SLL   _ _ x -> Just (pack $ extend x)
SRL   _ _ x -> Just (pack $ extend x)
SRA   _ _ x -> Just (pack $ extend x)
JAL   x    -> Just (pack $ extend x)
J     x    -> Just (pack $ extend x)
_       -> Nothing

```

The next part is little bit tricky: dispatching the ALU operation. Basic arithmetic operations can be dispatched directly. Notwithstanding, there are some special cases. For example, BEQ and BNE are dispatched to XOR because we want to use the zero flag of the xor result to check whether two operands are equal. Those memory related operations are dispatched to addition, because we want to add up the immediate value and the value of RS to get the target address. As for jump operations, we will pass the jump target via immediate value so we just use the OR operation (another operand will be 0).

`dispatch`

```

:: HiddenClockResetEnable dom
=> Signal dom Instruction
-> Signal dom ALUOperation
dispatch = fmap $
  \case
    ADD    _ _ _ -> ALUAdd True
    ADDI   _ _ _ -> ALUAdd True
    ADDU   _ _ _ -> ALUAdd False
    ADDIU  _ _ _ -> ALUAdd False
    SUB     _ _ _ -> ALUSub True
    SUBU    _ _ _ -> ALUSub False

```

```

AND    _ _ _ -> ALUAnd
ANDI   _ _ _ -> ALUAnd
NOR    _ _ _ -> ALUNor
OR     _ _ _ -> ALUOr
ORI    _ _ _ -> ALUOr
XOR    _ _ _ -> ALUXor
XORI   _ _ _ -> ALUXor
BEQ    _ _ _ -> ALUXor
BNE    _ _ _ -> ALUXor
SLT    _ _ _ -> ALUSet True
SLTI   _ _ _ -> ALUSet True
SLTU   _ _ _ -> ALUSet False
SLTIU  _ _ _ -> ALUSet False
LW     _ _ _ -> ALUAdd True
SW     _ _ _ -> ALUAdd True
SLL    _ _ _ -> ALUShiftL
SLLV   _ _ _ -> ALUShiftL
SRL    _ _ _ -> ALUShiftR False
SRLV   _ _ _ -> ALUShiftR False
SRA    _ _ _ -> ALUShiftR True
SRAV   _ _ _ -> ALUShiftR False
JR     _      -> ALUOr
J      _      -> ALUOr
JAL    _      -> ALUOr
_      _      -> ALUNone

```

Finally, we can finish the Control Unit interface:

```

{-# ANN controlUnit
  ( Synthesize
    { t_name = "ControlUnit"
    , t_inputs =
      [ PortName "CLOCK"
      , PortName "RESET"

```

```

        , PortName "ENABLE"
        , PortName "Instruction"
    ]
    , t_output =
        PortProduct "CTL"
        [ PortName "WRITE"
        , PortName "MEM"
        , PortName "BRANCH_FLAG"
        , PortName "ALU"
        , PortName "IMM"
        ]
    }
)
#-}

```

```
controlUnit
```

```

:: Clock System
-> Reset System
-> Enable System
-> Signal System Instruction           -- instruction
-> ( Signal System (Maybe (Unsigned 5)) -- write register
    , Signal System MemoryOperation    -- memory
    , Signal System BranchFlag         -- branch flag
    , Signal System ALUOperation       -- ALU control
    , Signal System (Maybe (BitVector 32)) -- immediate value
    )

```

```
controlUnit =
```

```

    exposeClockResetEnable $
    (,,,,)
    <$> writeRegister
    <*> memoryOperation
    <*> branchFlag

```

```

<*> dispatch
<*> immediateValue

```

As you can see, we just collect all the parts together, making no special change.

Decode Module Interface

Eventually, we complete our design of the Control Unit and get back to the Decode Module. Unfortunately, there is still one more thing to do: to decide registers to fetch. This is described in the following function:

```

registerPair
  :: HiddenClockResetEnable dom
  => Signal dom Instruction
  -> Signal dom (RegNo, RegNo)
registerPair = fmap $
  \case
    ADD    x y _ -> (x, y)
    ADDI   x _ _ -> (x, 0)
    ADDU   x y _ -> (x, y)
    ADDIU  x _ _ -> (x, 0)
    SUB    x y _ -> (x, y)
    SUBU   x y _ -> (x, y)
    AND    x y _ -> (x, y)
    ANDI   x _ _ -> (x, 0)
    NOR    x y _ -> (x, y)
    OR     x y _ -> (x, y)
    ORI    x _ _ -> (x, 0)
    XOR    x y _ -> (x, y)
    XORI   x _ _ -> (x, 0)
    BEQ    x y _ -> (x, y)
    BNE    x y _ -> (x, y)
    SLT    x y _ -> (x, y)
    SLTI   x _ _ -> (x, 0)

```

```

SLTU  x y _ -> (x, y)
SLTIU x _ _ -> (x, 0)
LW    x _ _ -> (x, 0)
SW    x y _ -> (x, y)
SLL   _ x _ -> (x, 0)
SRL   _ x _ -> (x, 0)
SRA   _ x _ -> (x, 0)
SLLV  x y _ -> (y, x)
SRLV  x y _ -> (y, x)
SRAV  x y _ -> (y, x)
NOP           -> (0, 0)
J        _      -> (0, 0)
JAL      _      -> (0, 0)
JR       x      -> (x, 0)

```

For those who requires immediate values, we just set the second register to zero and it will not be used. The first register of jump operations' will also be set to zero to make sure that OR operation will preserve the immediate value. Next, let us write a function to handle the state transition of the decode module: at the rising edge of the clock, this module will store the output as the new state and handle the decoding of the previous state.

```

type DecodeModuleState = (Instruction, Unsigned 32)

```

```

decodeModuleState

```

```

  :: (Instruction, Unsigned 32, Bool)

```

```

  -> State DecodeModuleState DecodeModuleState

```

```

decodeModuleState (inst, pc, stall) = do

```

```

  case stall of

```

```

    False -> do

```

```

      state <- get

```

```

      put (inst, pc)

```

```

      return state

```

```

    True -> do

```



```

    let res = (NOP, 0)
    put res
    return res

```

Special cases occur when the stall flag is set; it will then flush the output and state by setting the instruction to NOP.

```

{-# ANN decodeModule
  ( Synthesize
    { t_name = "DecodeModule"
    , t_inputs =
      [ PortName "CLOCK"
      , PortName "RESET"
      , PortName "ENABLE"
      , PortName "WRITE_REG"
      , PortName "STALL"
      , PortName "INSTRUCTION"
      , PortName "COUNTER"
      ]
    , t_output =
      PortProduct "DM"
      [ PortName "WRITE"
      , PortName "MEM"
      , PortName "BRANCH_FLAG"
      , PortName "ALU"
      , PortName "IMM"
      , PortName "RS"
      , PortName "RSV"
      , PortName "RT"
      , PortName "RTV"
      , PortName "COUNTER"
      ]
    }
  )

```

```

#-}
decodeModule
  :: Clock System
  -> Reset System
  -> Enable System
  -> Signal System ( Maybe (RegNo, Reg) ) -- write data
  -> Signal System Bool                  -- stall
  -> Signal System Instruction           -- instruction
  -> Signal System ( Unsigned 32 )       -- counter
  -> Signal System ( Maybe ( Unsigned 5 ) -- write register
    , MemoryOperation                  -- memory
    , BranchFlag                       -- branch flag
    , ALUOperation                     -- ALU control
    , Maybe ( BitVector 32 )           -- immediate value
    , RegNo                            -- rs
    , BitVector 32                     -- rs value
    , RegNo                            -- rd
    , BitVector 32                     -- rd value
    , Unsigned 32                      -- output counter
  )
decodeModule clk rst enable wdata stall inst counter =
  let stateMachine =
    exposeClockResetEnable $ asStateM decodeModuleState (NOP, 0)
  (rinst, pc) =
    unbundle (stateMachine clk rst enable $ bundle (inst, counter, stall))
  regDecoder =
    exposeClockResetEnable registerPair
  (rs, rt) =
    unbundle $ regDecoder clk rst enable rinst
  (w, m, b, a, i) =
    controlUnit clk rst enable rinst
  (rsv, rtv) =
    unbundle (registerFile clk rst enable $ bundle (rs, rt, wdata))

```

```
in bundle (w, m, b, a, i, rs, rsv, rt, rtv, pc)
```

Finally, we come up with the interface of the whole module. On clock rising, the module will first handle the state transition and get the instruction from its own register. Based on the instruction, it decides the register to read and generate the output in the control unit. One special part is that the register writing information is not from the state register and the module will handle the external write request at the same cycle. The read and write information are set to the register file. The indices of the register operands will also be passed to the next module, because they are useful to decide the forwarding information in ALU. On stalling, this module will just behaves as if the current instruction is NOP, but the writing requests are handled correctly.

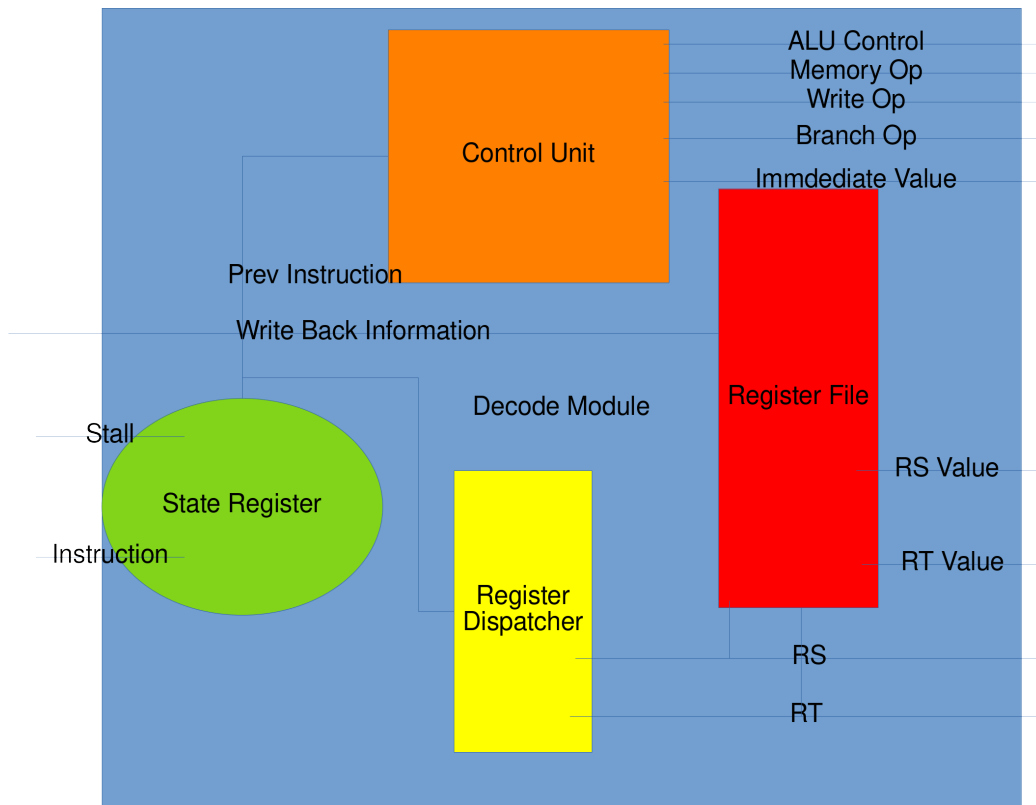


Figure 4.3: Decode Module Diagram

4.3.3 Arithmetic Module

Arithmetic Module is in charge of calculations. It will also check the branching condition and determine the target. Memory Operation cannot live without this module either, as the memory address is decided in this part.

Data Types

In order to make life easier, let us first define some data types:

- **MemoryOperation'**: This will be one of the output of this module. Previously, we have already defined **MemoryOperation**. This new data type is almost the same, but the write operation carries its writing value now as the value is ready after arithmetic operations.

```
data MemoryOperation'
  = MemNone'
  | MemLoad'
  | MemWrite' (BitVector 32)
deriving (Generic)
deriving (NFDataX)
deriving (Show)
```

- **ALUState**: just the same as the output product of the Decode Module.
- **ALUOutput**: a product type of the output of this module

```
type ALUOutput
  = ( Maybe (Unsigned 5) -- write register
    , MemoryOperation'   -- memory
    , BitVector 32        -- ALU result
    , Maybe (Unsigned 32) -- branch target
  )
```

ALU

ALU is a combinatorial circuit within the CPU that handles most of the arithmetic works. Here, our ALU follows the control command sent from the Control Unit and

apply the target operations on its operands. Apart from the arithmetic result, ALU will also generate some flags like zero, negative and overflow. Although only zero flag is useful for this book, we are going to implement all these flags.

For additions, overflow happens if the operands has the same sign bit while the sign bit differs in the result; for subtraction, overflow happens if the operands has different sign bit while the sign bit of the result is the same as the second operand. The following functions implement the detection of overflow:

```
addOverflow :: BitVector 32 -> BitVector 32 -> (BitVector 32, Bool)
addOverflow a b =
  let c = a + b
  in (c, (a ! 31) == (b ! 31) && (a ! 31) /= (c ! 31))
```

```
subOverflow :: BitVector 32 -> BitVector 32 -> (BitVector 32, Bool)
subOverflow a b =
  let c = a - b
  in (c, (a ! 31) /= (b ! 31) && (c ! 31) == (b ! 31))
```

The interface of ALU is described as the following:

```
type ALUResult
= ( BitVector 32 -- Arithmetic Result
  , Bool -- Overflow Result
  , Bool -- Zero Flag
  , Bool -- Negative Flag
  )
```

```
{-# ANN arithmeticUnit
  ( Synthesize
    { t_name = "ArithmeticUnit"
    , t_inputs =
      [ PortName "OPERATION"
      , PortName "OPERAND_1"
      , PortName "OPERAND_2"
```

```

    ]
    , t_output =
        PortProduct "ALU"
        [ PortName "RESULT"
        , PortName "OVERFLOW"
        , PortName "ZERO"
        , PortName "NEG"
        ]
    }
)
#-}

```

```

arithmeticUnit :: ALUOperation -> BitVector 32 -> BitVector 32 -> ALUResult
arithmeticUnit op opr0 opr1 = (res, o, z, n)
  where
    (res, o, n) = arithmeticUnit' op
    z = res == 0

```

The following part is the inner implementation of each operation:

1. Addition and Subtraction: signed and unsigned addition change the bits in the same way, but only signed operation may overflow.

```

arithmeticUnit' (ALUAdd flag) =
  let (res, overflow) = addOverflow opr0 opr1
  in (res, overflow && flag, bitToBool $ res ! 31)
arithmeticUnit' (ALUSub flag) =
  let (res, overflow) = subOverflow opr0 opr1
  in (res, overflow && flag, bitToBool $ res ! 31)

```

2. And, Or, Xor, Nor are simple bitwise operations:

```

arithmeticUnit' ALUAnd = (opr0 .&. opr1, False, False)
arithmeticUnit' ALUNor = (complement $ opr0 .|. opr1, False, False)
arithmeticUnit' ALUOr  = (opr0 .|. opr1, False, False)
arithmeticUnit' ALUXor  = (opr0 `xor` opr1, False, False)

```

3. Set-on-less-than needs to distinguish the sign information:

```
arithmeticUnit' (ALUSet flag) =
  let result =
    boolToBV $
      if flag
      then (unpack opr0 :: Signed 32) < (unpack opr1 :: Signed 32)
      else (unpack opr0 :: Unsigned 32) < (unpack opr1 :: Unsigned 32)
  in (result, False, False)
```

4. Shift operations can use the predefined functions in Clash but we need to check the sign information for the right shifting:

```
arithmeticUnit' ALUShiftL =
  (opr0 `shiftL` (unpack $ extend opr1), False, False)
arithmeticUnit' (ALUShiftR True) =
  ( pack $ (unpack opr0 :: Signed 32) `shiftR` (unpack $ extend opr1)
  , False
  , False)
arithmeticUnit' (ALUShiftR False) =
  (opr0 `unsafeShiftR` (unpack $ extend opr1), False, False)
```

5. Non-Operation does nothing:

```
arithmeticUnit' ALUNone = (0, False, False)
```

Forward Unit

As we have mentioned, we must handle some data hazards using the forward strategy. This is done in the Forward Unit of Arithmetic Module.

```
type ForwardInfo = Maybe (Unsigned 5, BitVector 32)
```

```
forwardUnit'
  :: HiddenClockResetEnable dom
  => Signal dom ForwardInfo    -- mem register
```

```

-> Signal dom ForwardInfo    -- write-back register
-> Signal dom (Unsigned 5)    -- rs
-> Signal dom (Unsigned 5)    -- rt
-> Signal dom (Maybe (BitVector 32), Maybe (BitVector 32))
forwardUnit' a b c d = forwarding <$> a <*> b <*> c <*> d
where
  forwarding exec load rs rt =
    let rs' =
      case (exec, load) of
        (Just (no, res), _)
          | no == rs -> Just res
        (_, Just (no, res))
          | no == rs -> Just res
        _          -> Nothing
    rt' =
      case (exec, load) of
        (Just (no, res), _)
          | no == rt -> Just res
        (_, Just (no, res))
          | no == rt -> Just res
        _          -> Nothing
    in (rs', rt')

{-# ANN forwardUnit
  ( Synthesize
    { t_name = "ForwardUnit"
    , t_inputs =
      [ PortName "CLOCK"
      , PortName "RESET"
      , PortName "ENABLE"
      , PortName "FORWARD_A"
      , PortName "FORWARD_B"
      , PortName "RS"

```



```

        , PortName "RT"
      ]
    , t_output =
      PortProduct "FW"
      [ PortName "OVERRIDE_RS"
      , PortName "OVERRIDE_RT"
      ]
    }
  )
#-}

forwardUnit
  :: Clock System
  -> Reset System
  -> Enable System
  -> Signal System ForwardInfo -- memm register
  -> Signal System ForwardInfo -- load register
  -> Signal System (Unsigned 5) -- rs
  -> Signal System (Unsigned 5) -- rt
  -> Signal System (Maybe (BitVector 32), Maybe (BitVector 32))
forwardUnit = exposeClockResetEnable forwardUnit'

```

This unit has two ports to accept forward data and two ports to get the register numbers. The forward data comes from two successor stages: memory stage and write-back stage; whenever there is a write-back request being issued in these two stages, it will be sent to the Forward Unit. If the write-back register happens to be the same as the operand register in ALU, then the data hazard is found. What if memory stage and write-back stage changes the same register? Notice that the changes in the write-back are already applied to the memory part by a previous forward, thus we can just accept the forward in the write-back part to get the latest value of the register. This handling method is expressed in the pattern matching order at the code above.

State Machine

The state transition of this module is also simple: when stalling happens, it just flushes away the current operations and states; otherwise, pulls out the operations from the current state and stores the new state.

```

arithmeticModuleState' :: (ALUState, Bool) -> State ALUState ALUState
arithmeticModuleState' (_, True) = do
    let res
        = (Nothing, MemNone, NoBranch, ALUNone, Nothing, 0, 0, 0, 0, 0)
    put res
    return res
arithmeticModuleState' (state, False) = do
    res <- get
    put state
    return res

arithmeticModuleState
    :: HiddenClockResetEnable dom
    => Signal dom (ALUState, Bool)
    -> Signal dom ALUState
arithmeticModuleState =
    asStateM arithmeticModuleState'
    (Nothing, MemNone, NoBranch, ALUNone, Nothing, 0, 0, 0, 0, 0)

```

Arithmetic Module Interface

Arithmetic Module has the most complicated logic relation. It must check a lot of things and do the arithmetic operations.

```

1 {-# ANN arithmeticModule
2     ( Synthesize
3       { t_name = "AithmeticModule"
4         , t_inputs =
5           [ PortName "CLOCK"

```

```

6         , PortName "RESET"
7         , PortName "ENABLE"
8         , PortName "FW_0"
9         , PortName "FW_1"
10        , PortName "STALL"
11        , PortProduct "AM"
12          [ PortName "WRITE"
13            , PortName "MEM"
14            , PortName "BRANCH_FLAG"
15            , PortName "ALU"
16            , PortName "IMM"
17            , PortName "RS"
18            , PortName "RSV"
19            , PortName "RT"
20            , PortName "RTV"
21            , PortName "COUNTER"
22          ]
23      ]
24      , t_output =
25        PortProduct "AM"
26          [ PortName "WRITE_REG"
27            , PortName "MEM_OP"
28            , PortName "RESULT"
29            , PortName "BRANCH_TARGET"
30          ]
31      }
32  )
33  #-}
34  arithmeticModule
35      :: Clock System
36      -> Reset System
37      -> Enable System

```

```

38   -> Signal System ForwardInfo -- last
39   -> Signal System ForwardInfo -- last last
40   -> Signal System Bool
41   -> Signal System ALUState    -- input
42   -> Signal System ALUOutput  -- output
43   arithmeticModule clk rst enable last last' stall input =
44     let (write, mem, branch, alu, imm, rs, rsv, rt, rtv, counter) =
45         unbundle $
46         ((exposeClockResetEnable arithmeticModuleState) clk rst
47          ↪ enable) $
48         bundle (input, stall)
49
50     (check0, check1) = unbundle $ forwardUnit clk rst enable
51     ↪ last last' rs rt
52
53     unwrap (Just a) = a
54     unwrap _ = 0
55
56     rsv' = (<|>) <$> check0 <*> (pure <$> rsv)
57     rtv0 = ((<|>) <$> check1 <*> (pure <$> rtv))
58     rtv' = (<|>) <$> imm <*> rtv0
59
60     memSolver MemWrite value = MemWrite' value
61     memSolver MemLoad _ = MemLoad'
62     memSolver _ _ = MemNone'
63
64     mem' = memSolver <$> mem <*> (unwrap <$> rtv0)
65     (res, _, z, _) =
66         unbundle $
67         arithmeticUnit <$> alu <*> (unwrap <$> rsv') <*> (unwrap
68         ↪ <$> rtv')

```

```

66     check_branch True (BranchEQ delta) pc _ = Just (pc + unpack
        ↪ delta)
67     check_branch False (BranchNE delta) pc _ = Just (pc +
        ↪ unpack delta)
68     check_branch _ Jump _ (Just i) = Just (unpack i)
69     check_branch _ _ _ _ = Nothing
70
71     branch' = check_branch <$> z <*> branch <*> counter <*> imm
72 in bundle (write, mem', res, branch')
```

Indeed, this module has a lot of things to do.

1. Line 44 to line 47 finish the state transition.
2. Line 48 checks the forwarding information.
3. Line 53 to line 55 determines different register values. For RS, if there is a forward value, we will override the value to the latest version; For RT, we do mostly the same, but for the Arithmetic Unit, if there is an immediate value, the final value should be `rtv'`; but RT can also be used in the output of memory operation flags for the next stage, in that case, we will use `rt0`.
4. Line 57-61 decides the memory flags and values.
5. Line 62-64 represents the ALU operations.
6. Line 66-71 checks branching/jumping and calculates the target.

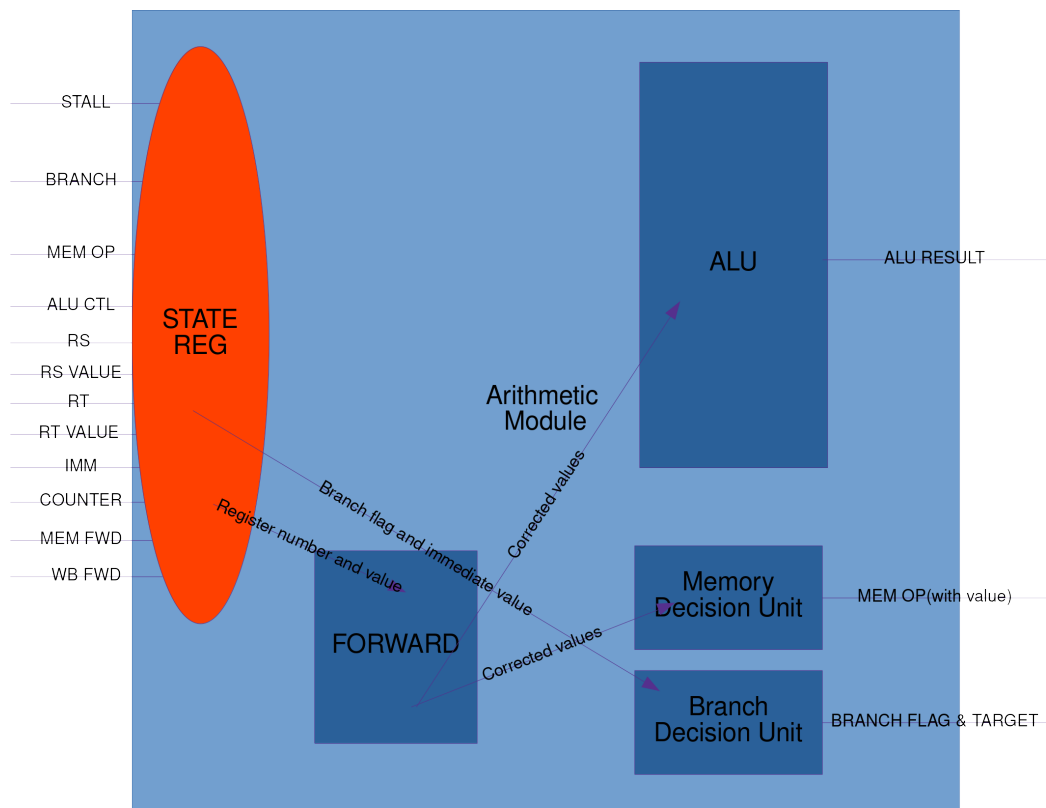


Figure 4.4: Arithmetic Module Diagram)

4.3.4 Memory Module

Main Memory Block

The main memory also uses the BRAM structure. It is implemented in the following:

```

1  {-# ANN mainRAM
2      ( Synthesize
3          { t_name = "MainMemory"
4            , t_inputs =
5              [ PortName "CLOCK"
6                , PortName "RESET"
7                , PortName "ENABLE"
8                , PortName "FETCH_ADDRESS"
9                , PortName "EDIT_SERIAL"

```

```

10         ]
11         , t_output = PortName "DATA"
12     }
13 )
14 #-}
15
16 mainRAM
17   :: Clock System
18   -> Reset System
19   -> Enable System
20   -> Signal System MemAddr
21   -> Signal System (Maybe (MemAddr, MemoryBlock))
22   -> Signal System MemoryBlock
23 mainRAM = exposeClockResetEnable $ blockRam (replicate d512 0)

```

The block memory adapts the default write-after-read strategy and it takes two inputs: the read address and the write information.

Memory State

Again, the state transition is mainly about exchange the output data from the previous module.

```

1 memState' :: (ALUOutput, Bool) -> State ALUOutput ALUOutput
2 memState' (_, True) = do
3   let res = (Nothing, MemNone', 0, Nothing)
4   put res
5   return res
6 memState' (state, _) = do
7   res <- get
8   put state
9   return res
10
11 memState
12   :: HiddenClockResetEnable dom

```

```

13     => Signal dom (ALUOutput, Bool)
14     -> Signal dom ALUOutput
15 memState = asStateM memState' (Nothing, MemNone', 0, Nothing)

```

Memory Module Interface

```

1  type MemOutput
2    = ( Maybe (Unsigned 32) -- branch target
3        , Maybe (RegNo, Reg) -- write pair
4        )
5
6  {-# ANN memoryModule
7    ( Synthesize
8      { t_name = "MemoryModule"
9        , t_inputs =
10          [ PortName "CLOCK"
11            , PortName "RESET"
12            , PortName "ENABLE"
13            , PortProduct "MMI"
14              [ PortName "WRITE_REG"
15                , PortName "MEM_OP"
16                , PortName "RESULT"
17                , PortName "BRANCH_TARGET"
18              ]
19            , PortName "STALL"
20          ]
21        , t_output =
22          PortProduct "MMO"
23            [ PortName "BRANCH"
24              , PortName "WRITE_PAIR"
25            ]
26        }
27    )

```



```

28  #-}
29  memoryModule
30      :: Clock System
31      -> Reset System
32      -> Enable System
33      -> Signal System ALUOutput
34      -> Signal System Bool
35      -> Signal System MemOutput
36  memoryModule clk rst enable aluOut stall =
37      let stateMachine =
38          (exposeClockResetEnable memState) clk rst enable
39          (writeReg, memOpPrev, aluRes, br) =
40          unbundle $ stateMachine $ bundle (aluOut, stall)
41
42      writeInfoSolver _ _ True _ = Nothing
43      writeInfoSolver address (MemWrite' v) _ Nothing =
44          Just (unpack address, v)
45      writeInfoSolver _ _ _ _ = Nothing
46
47      thirdData (_, _, x, _) = x `unsafeShiftR` 2
48      secondData (_, x, _, _) = x
49
50      aluRes' = thirdData <$> aluOut
51      memOp   = secondData <$> aluOut
52
53      writeInfo =
54          writeInfoSolver
55              <$> aluRes' <*> memOp <*> stall <*> br
56
57      memData =
58          mainRAM clk rst enable (unpack <$> aluRes') writeInfo
59

```

```

60     writePair (Just no) _ MemLoad' res = Just (no, res)
61     writePair (Just no) res _ _ = Just (no, res)
62     writePair _ _ _ _ = Nothing
63
64     writePair' =
65         writePair
66             <$> writeReg <*> aluRes <*> memOpPrev <*> memData
67
68     check True _ = (Nothing, Nothing)
69     check _ x = x
70 in check <$> stall <*> bundle (br, writePair')
```

Please notice that `aluRes'` and `memOp` comes from the Arithmetic Module at the same cycle, this makes sure that in the next cycle, the BRAM will output the correct data.

The memory read/write information is derived from the instant input while the register write information comes from the state register; this prevents possible oscillations and fits well with the BRAM delay effects.

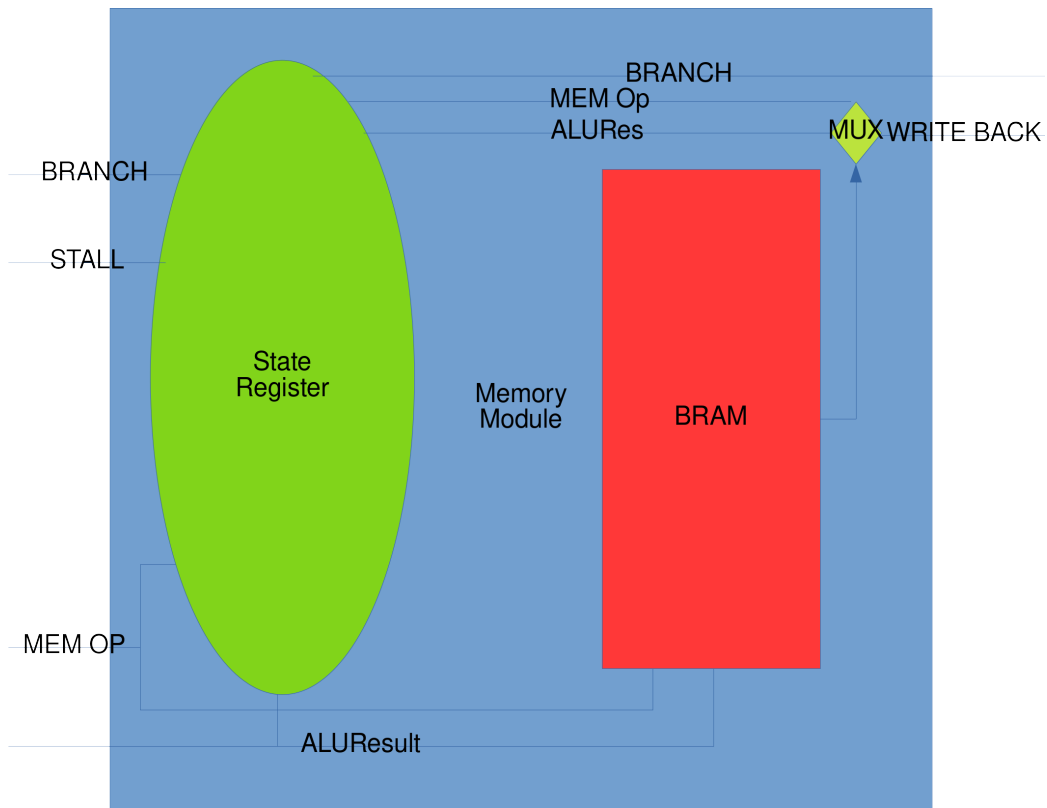


Figure 4.5: Memory Module Diagram

4.3.5 Write-Back Module

Write-Back is the simplest module: just a buffer before editing the data.

State Register

The state of this module is simply a register. In this part, we do not need to care about stalling.

```

1 writeRegister
2   :: HiddenClockResetEnable dom
3   => Signal dom (Maybe (Unsigned 32), Maybe (RegNo, Reg))
4   -> Signal dom (Maybe (Unsigned 32), Maybe (RegNo, Reg))
5 writeRegister = register (Nothing, Nothing)

```

Write-Back Module Interface

The interface is nothing more than a wrapper around the register:

```

1  {-# ANN writeBack
2    ( Synthesize
3      { t_name = "WriteBack"
4        , t_inputs =
5          [ PortName "CLOCK"
6            , PortName "RESET"
7            , PortName "ENABLE"
8            , PortName "BRANCH"
9            , PortName "WRITE_PAIR"
10           ]
11        , t_output =
12          PortProduct "WB"
13            [ PortName "BRANCH"
14              , PortName "WRITE_PAIR"
15            ]
16        }
17      )
18  #-}
19  writeBack
20    :: Clock System
21    -> Reset System
22    -> Enable System
23    -> Signal System (Maybe (Unsigned 32))
24    -> Signal System (Maybe (RegNo, Reg))
25    -> Signal System (Maybe (Unsigned 32), Maybe (RegNo, Reg))
26  writeBack clk rs en br wb = writeRegister' $ bundle (br, wb)
27  where
28    writeRegister' = (exposeClockResetEnable writeRegister) clk
                      ↪ rs en

```

Advanced Topics

Constraint Trick with `-XPartialTypeSignatures`

This part is provided by dramforever.

Tired of writing all the

`HiddenClockResetEnable` dom, `NFDataX` a, `Num` a, `Eq` a

constraints? Just enable `-XPartialTypeSignatures` and `-Wno-partial-type-signatures` then replace all your constraints with a `_`.

Before:

```
foo
  ::
  ( HiddenClockResetEnable dom
    , Num a, Default a, NFDataX a
    )
  => a -> a -> a
  -> Signal dom a
  -> Signal dom a
```

After:

```
foo
  :: _
  => a -> a -> a
  -> Signal dom a
  -> Signal dom a
```

Synthesize Annotation Trick

This part is provided by `dramforever`.

Have you ever found yourself writing **Synthesize** annotations like this:

```
{-# ANN f Synthesize
  { t_name    = "f"
    , t_inputs = [ PortName "a"
                  , PortProduct "" [ PortName "b", PortName "c" ] ]
    , t_output = PortProduct "res" [PortName "q"]
  } #-}
```

And thought, maybe there's a less verbose way. Well, look no further than **-XOverloadedLists** and **-XOverloadedStrings**

```
instance IsString PortName where
```

```
  fromString = PortName
```

```
instance IsList PortName where
```

```
  type Item PortName = PortName
```

```
  fromList = PortProduct ""
```

```
  toList = error "toList for PortName is not implemented"
```

Now we can just write:

```
{-# ANN f Synthesize
  { t_name    = "f"
    , t_inputs = [ "a", [ "b", "c" ] ]
    , t_output = PortProduct "res" [ "q" ]
  } #-}
```

Isn't that just like, way more readable?

Note that these are indeed orphans so you might want to put `{-# OPTIONS -Wno-orphans #-}` in whatever file you are defining these instances.

Applicative Do-Block

Sometimes, if you are describing some combinatorial logic within signal, it may be useful to use Applicative Do-Block. You need to enable the ApplicativeDo language extension:

```
test :: Signal System Bool
      -> Signal System Bool
      -> Signal System Bool
      -> Signal System (Bool, Bool)
test signalA signalB signalC = do
  a <- signalA
  b <- signalB
  c <- signalC
  return (a && b, b || c)
```

This is transformed into the following thing:

```
(\a b c -> (a && b, b || c)) <$> a <*> b <*> c
```

Be aware that Signal is not a Monad Instance. Even though you can use some basic logic structures, these do blocks are highly limited: you can not use let-binding or apply some complicated functions at the middle way.

GtkWave

In the test bench, we have already shown the functions to dump VCD files. The dumped data can be used in a software called GtkWave to provide visualization of the circuit state.

Invoke Icarus Verilog simulation executable and you will get the monitor output together with a VCD file.

You can open the file with GtkWave and get a screen like this:

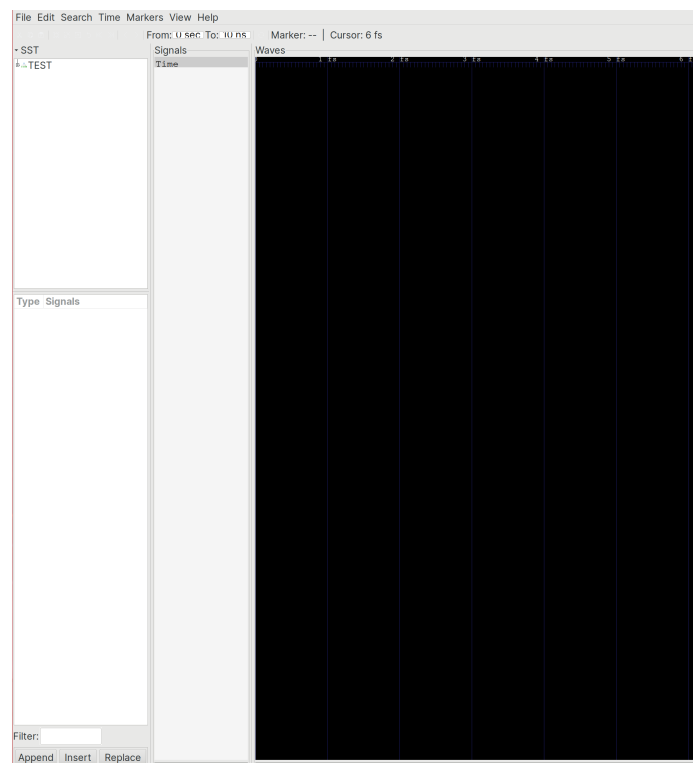


Figure 6: GtkWave Screenshot

Using the tree view on the left-hand side, you can choose the signals to display, for example, the following screenshot shows the graph of our toy CPU under several clock cycles:

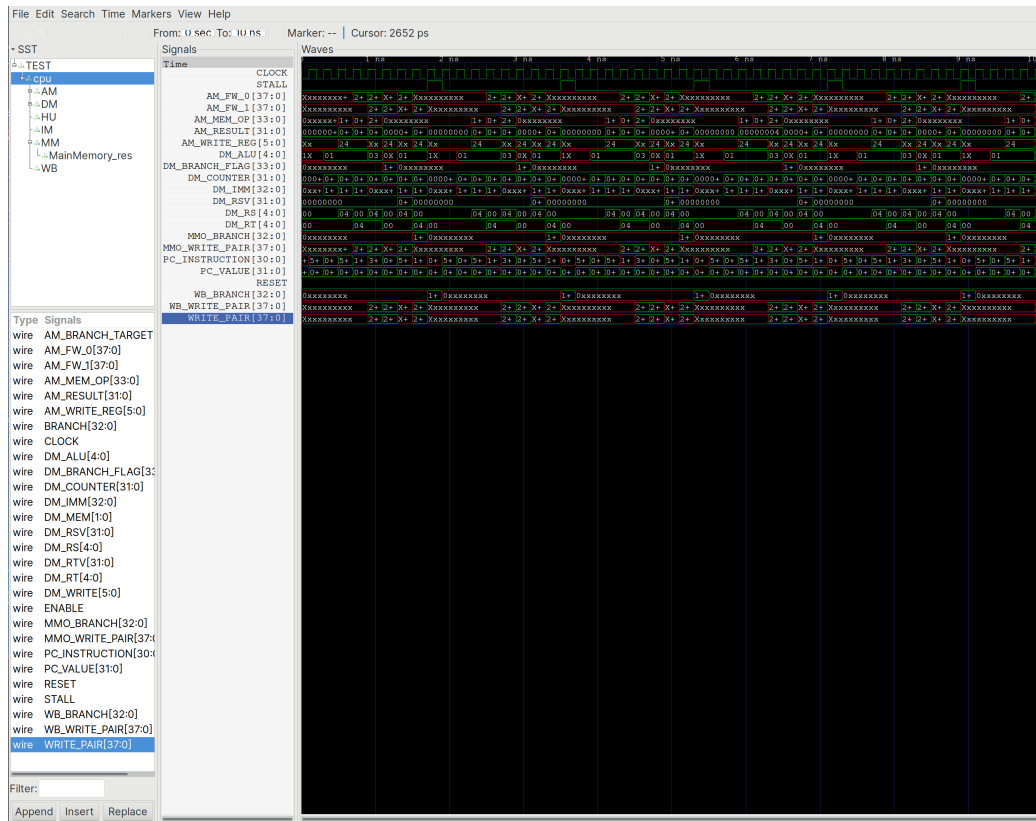


Figure 7: GtkWave Display

If you want to use GtkWave to debug the circuits, it is recommended to store the configuration so that you do not need to select the signals to display each time you open the program. In the command-line, you can invoke the program by something like this:

```
gtkwave test.vcd config.gtkw
```

Sample Output

Let us create a nonsense assembly file which involves stalling and forwarding:

```
1  lw $t0, 4($zero)
2  addi $t0, $t0, 1
3  lw $t1, 8($zero)
4  addi $t1, $t1, 1
5  sw  $t0, 4($zero)
6  sw  $t1, 8($zero)
7  sub  $t0, $t0, $t0
8  xor  $t1, $t1, $t1
9  beq  $zero, $zero, 0
10 addi $t0, $t0, 1
11 addi $t1, $t1, 1
12 j 114514
```

This file is translated into the following binary form:

```
1  1000110000000100000000000000000100
2  001000010000100000000000000000001
3  1000110000000100100000000000001000
4  001000010010100100000000000000001
5  101011000000010000000000000000100
6  1010110000000100100000000000001000
7  00000001000010000100000000100010
8  00000001001010010100100000100110
9  00010000000000001111111111110111
```

```
10 0010000100001000000000000000000001
11 0010000100101001000000000000000001
12 000010000000000000110111111010100
```

The output is shown in the following files:

1. output.txt:

<https://yifan.cowtransfer.com/s/166018e0ad394b>

2. output.vcd:

<https://yifan.cowtransfer.com/s/e10ff6693c3149>

Acknowledgments

This book cannot be finished without the help from some excellent people. We must give our acknowledgments to them.

Thanks **Prof. Ye-Ching Chung** <ychung@cuhk.edu.cn> for his lectures and guidance on microprocessor architectures and instruction sets.

Thanks **dramforever** <dramforever@live.com> for the template repository and several important corrections of our technical details, mainly on the installation part and the descriptions of several Clash functions.

Thanks **Mickey Ma** <mickeyma@cuhk.edu.cn> for her kind and patient advice on the details of the CPU design.

Thanks **Outvi V** <i@outv.im> for providing some materials of GtkWave and good suggestions on hazards handling.

Thanks **ice1000** <ice1000kotlin@foxmail.com> for the suggestions on Haskell code-style and logic structures.